NUR - Hand-in Exercise 1

Martina Cacciola (s4170814)

March 7, 2024

Abstract

In this document, I present the results for the Hand-in Exercise 1 for the course Numerical Recipes for Astrophysics.

1 Poisson distribution

The exercise is done in the script

```
#!/usr/bin/env python
  import numpy as np
  from math import exp, log10
  This script containes the code for Exercise 1.
  def poisson_prob(l, k):
       Calculates the Poisson probability for a given lambda and k
      At each passage, we work with 32 bit floats and in log space to avoid overflow.
      By taking the logarithm, we transform multiplications into additions and divisions
13
      into subtractions.
      This helps keep the intermediate results within a manageable range.
      We then take the exponent of the final result, ensuring both numerical stability and
        correct results.
17
      Inputs:
      l: lambda value
18
      k: k value
19
      Output:
20
21
       Poisson probability for a given lambda and k
22
23
      We divide the calculation into:
       1. Calculate lamba<sup>k</sup> in log space
24
       2. Calculate e^-lambda
25
       3. Calculate k! and the cumulative sum of log(k!) for all k from 2.
      The reason for starting from 2 is that the first two values of k! are 0 and 1, not
       useful for the calculation.
      The cumulative sum is useful for next calculation
       4. Calculate the Poisson probability in log space and convert it back to normal
29
      space.
30
31
32
       k_{\text{values}} = \text{np.float32}(\text{np.arange}(k + 1))
33
34
      # 1. log(lamba^k)
35
      lam_k log = np. float 32 (k_values * log 10(l))
36
37
      # 2. e^-lambda
      e_neg_lambda = np.float32(log10(exp(-1)))
38
      # 3. Factorial calculation: log(k!) and cumulative sum of log(k!) for all k from 2
39
40
      log_fact_k = np. float 32 (np. log 10 (np. arange (2, k + 1)))
       tot_log_fact = np.float32(np.cumsum(log_fact_k))
41
      # 4. Poisson probability in log space and conversion
```

```
\log_p k = np. float32 (lam_k log [2:] + e_neg_lambda - tot_log_fact)
       # Terms k=0,1 are calculated separately to maintain accuracy
       # The remaining k are obtained by exponentiating the logarithmic Poisson
45
       probabilities
       poisson\_distribution = np.float32(np.concatenate([[(l**k) * exp(-l)] for k in
       k_values[:2]] + [np.power(10, log_p_k)]))
47
       return poisson_distribution
48
49
  # Test the function for the given values
  parameter_values = [(1, 0), (5, 10), (3, 21), (2.6, 40), (101, 200)]
51
  # Save Poisson probability results to a text file
  output_data = []
54
  for lambda_val, k_val in parameter_values:
       poisson_distribution = poisson_prob(lambda_val, k_val)
       output\_data.append\left(\left[lambda\_val\;,\;\;k\_val\;,\;\;poisson\_distribution\left[\left.int\left(k\_val\right)\right]\right]\right)
57
  np.savetxt(\ 'poisson\_output.txt',\ np.array(output\_data),\ header='Lambda\ 'tK\ tPoissonProbability',\ delimiter='\ 't')
```

poisson.py

. The necessary explanations are in the comments of the code. The output produced for the values required is in

```
# Lambda
      \mathbf{K}
        Poisson Probability
3.678794503211975098e-01
1.813277788460254669\,\mathrm{e}{-02}
1.019341815816643404e\!-\!11
2.600000000000000089e+00
               3.615135162986033128e - 33
1.269425385826567513\mathrm{e}{-18}
```

poisson_output.txt

.

2 Vandermonde matrix

In this section, we add the comments to the plots produced by the script given by:

```
#!/usr/bin/env python
  import numpy as np
  import matplotlib.pyplot as plt
  import timeit
  This script contains the code for Exercise 2.
  Point a):
  We implement the LU decomposition using Crout's algorithm to solve the Vandermonde
13
  The problem is divided into:
14
  1. LU decomposition: the matrix A is decomposed into a lower triangular matrix L and an
      upper triangular matrix U.
  Firstly we initialize L and U as zero matrices of the same size as A.
  Setting alpha_ii is done implicitly in passage 1:
  since L[i\,,\,\,:i\,] and U[:i\,,\,\,i\,] are both zero vectors at this point this dot product is zero
       , and so U[i, i] is just A[i, i].
  Then we update the elements of L and U loop over the columns.
20
  2. Forward substitution: we solve Ly = b for y.
  We initialize y as a zero vector of the same size as b.
25 Then we solve the system iteratively for each element of y starting from first element.
```

```
3. Backward substitution: we solve Ux = y for x.
  We initialize x as a zero vector of the same size as y.
  Then we solve the system iteratively for each element of x starting from the last
29
        element
30
31
   # LU Decomposition
32
   \# beta_ij corresponds to U[i, j], and alpha_ij corresponds to L[i, j]
33
   def lu_decomposition(A):
35
        n = len(A)
36
        L = np.zeros((n, n))
37
        U = np.zeros((n, n))
38
39
        # Loop over the columns
40
        for i in range(n): #Â Update the upper and lower triangular matrices
41
              \begin{array}{l} U[\,i\;,\;\;i\;:] \,=\, A[\,i\;,\;\;i\;:] \,-\, np.\, dot\, (L[\,i\;,\;\;i\;]\;,\;\; U[:i\;,\;\;i\;:]) \,\,\#\,\, Passage \,\, 1 \\ L[\,i\::\;,\;\;i\;] \,=\, (A[\,i\::\;,\;\;i\;] \,-\, np.\, dot\, (L[\,i\::\;,\;\;i\;]\;,\;\; U[:i\;,\;\;i\;])) \,\,/\,\, U[\,i\;,\;\;i\;] \,\,\#\,\, Passage \,\, 2 \\ \end{array} 
42
43
44
        return L, U
45
46
   # Forward substitution
47
   def forward_substitution(L, b):
48
        n = L.shape[0]
        y = np.zeros_like(b, dtype=np.double)
50
        # Solve Ly = b for y y[0] = b[0] / L[0, 0]
5.1
        for i in range(1, n):
53
             y[i] = (b[i] - np.dot(L[i, :i], y[:i])) / L[i, i] # Passage 3
54
        return y
56
   # Backward substitution
57
58
   def backward_substitution(U, y):
        n = U.shape[0]
50
        x = np.zeros_like(y, dtype=np.double)
        # Solve Ux = y for x
 x[-1] = y[-1] / U[-1, -1]
61
62
        for i in range (n-2, -1, -1):
63
             x[i] = (y[i] - np.dot(U[i, i:], x[i:])) / U[i, i] # Passage 4
64
65
        return x
66
   # Load the data
67
   \mathtt{data} = \mathtt{np.genfromtxt} \, (\texttt{"./Vandermonde.txt"}, \; \mathtt{comments} = \texttt{'\#'}, \; \mathtt{dtype} = \mathtt{np.float64})
   x \,=\, data\,[\,:\,,\quad 0\,]
69
  y = data[:, 1]
70
71
   # Create the Vandermonde matrix
72
   \begin{array}{ll} \textbf{def} & \texttt{create\_vandermonde} \, (\texttt{x} \,, \, \, \, N\!\!\!=\!\! None) \, ; \end{array}
73
        if N is None: # If N is not provided, use the length of x
74
             N = len(x)
        V = np.zeros((N, N))
76
77
        for i in range(N):
              for j in range(N):
78
                   V[i, j] = x[i] ** j
79
        return V
80
   V = create_vandermonde(x)
   # LU decomposition using Crout's algorithm
85
  L, U = lu\_decomposition(V)
   # Solve the system
   y_forward = forward_substitution(L, y)
   c = backward_substitution(U, y_forward)
   # Save the coefficients to a text file
91
92
  np.savetxt('coefficients.txt', c)
94 # Plot the 19-th degree polynomial evaluated at 1000 equally-spaced points
```

```
# Along with the original data points
   # This plot is just to have another look at the region of interest (it is the same as
       the next one)
   x_{\text{values}} = \text{np.linspace}(0, 100, 1000)
97
   y_values = np.zeros_like(x_values)
   for i in range(len(c)):
        y_values += c[i] * x_values ** i
   plt.plot(x_values, y_values, label="19-th degree polynomial")
   plt.ylim(-1000, 6000)
   plt.scatter(x, y, color='red', label='Data points')
   plt.legend()
104
   plt.savefig('./plots/polynomial_plot.png')
106
   xx=np.linspace(x[0],x[-1],1001) #x values to interpolate at
108
   # Calculate interpolated y values
   yya = np.array\left(\left[np.sum(c * xi ** np.arange(len(c))\right) for xi in xx\right]\right)
   ya = np.array([np.sum(c * xi ** np.arange(len(c))) for xi in x])
112
   # Plot 2a
114
   fig=plt.figure()
   gs=fig.add\_gridspec(2,hspace=0,height\_ratios=[2.0,1.0])
   axs=gs.subplots(sharex=True,sharey=False)
   axs[0].plot(x,y,marker='o',linewidth=0)
117
   plt.xlim(-1,101)
118
   axs[0].set_ylim(-400,400)
   axs[0].set_ylabel('$y$')
120
   #axs[1].set_ylim(1e-16,1e1)
121
   axs[1].set_ylim(1e-20,1e1)
   axs[1]. set_ylabel('$|y-y_i|$')
axs[1]. set_xlabel('$x$')
124
   axs[1]. set_yscale(''log')
   line,=axs[0].plot(xx,yya,color='orange')
126
   line.set_label('Via LU decomposition')
127
   axs [0]. legend (frameon=False, loc="lower left")
128
   axs\,[\,1\,]\,.\,plot\,(\,x\,, \\ \textcolor{red}{abs}\,(\,y-ya\,)\,\,,\\ \textcolor{red}{color=\,'\,orange\,\,'}\,)
129
   plt.savefig('./plots/my_vandermonde_sol_2a.png',dpi=600)
130
132
   Point b):
134
   We implement Neville's algorithm to interpolate the data points.
135
   1) We start by identifying the M tabulated points closest to the point x0
136
   To do so, we sort the points by their distance to x0 (using the function argsort) and
        select the M closest points.
      We set the initial p_i = y_i for i in 0 to M.
138
   3) We loop over each order of the interpolation from 1 to M\!-\!1
   For each order, we loop over each interval of the current order from 0 to M-k-1
   and update the p value for the interval, overwriting the previous orders' values.
   4) We return the first element of p, which is the interpolated value.
143
144
145
146
147
   def argsort (seq):
       # Pair each element in the sequence with its index
148
       indexed_seq = [(val, i) for i, val in enumerate(seq)]
149
150
       # Sort the indexed sequence
        sorted_indexed_seq = sorted(indexed_seq)
151
       # Extract the indices from the sorted indexed sequence
        sorted_indices = [i for val, i in sorted_indexed_seq]
        return sorted_indices
154
   def neville(x, y, x0, M):
157
158
        Parameters:
159
       x: x-coordinates of the data points
160
       y: y-coordinates of the data points
161
       x0: x-coordinate to interpolate at
162
```

```
M: number of points to use for interpolation
163
        Returns:
164
        p[0]: interpolated value at x0
165
166
        # Sort the points by their distance to x0 and select the M closest points
167
        sorted_indices = argsort(np.abs(x - x0))
168
        x = x[sorted\_indices[:M]]
169
170
        y = y[sorted_indices[:M]]
        \# Set p = y
        p = np.copy(y)
174
        # Loop over each order of the interpolation
175
        for k in range(1, M):
            # Loop over each interval of the current order
177
             for i in range (M-k):
178
                 # Update the p value for the interval
179
                 p[i] = ((x0 - x[i+k])*p[i] + (x[i] - x0)*p[i+1]) / (x[i] - x[i+k])
180
        return p[0] # Return the first element which is the interpolated value
181
182
183
   # Values for interpolation
184
   xx = np.linspace(0, 100, 1000)
185
186
   # Interpolation using Neville's algorithm
187
   yyb = np.array([neville(x, y, xi, 20) for xi in xx])

yb = np.array([neville(x, y, xi, 20) for xi in x])
189
190
   # Plotting
191
   fig = plt.figure()
192
   gs = fig.add_gridspec(2, hspace=0, height_ratios=[2.0, 1.0])
193
   axs = gs.subplots(sharex=True, sharey=False)
194
   axs[0].plot(x, y, marker='o', linewidth=0)
195
196
   plt.xlim(-1, 101)
   axs[0].set_ylim(-400, 400)
197
   axs[0].set_ylabel('$y$')
198
   axs[1].set_ylim(1e-30,1e1)
199
   axs[1]. set_ylabel('$|y-y_i|$')
axs[1]. set_xlabel('$x$')
axs[1]. set_yscale('log')
200
201
202
   line , = axs[0].plot(xx, yyb, linestyle='dashed', color='green')
line.set_label("Via Neville's algorithm")
203
   axs[0].legend(frameon=False, loc="lower left")
axs[1].plot(x, abs(y - yb), linestyle='dashed', color='green')
205
206
   plt.savefig('./plots/my_vandermonde_sol_2b.png', dpi=600)
208
209
   Point c):
210
   Iterative version of LU decomposition.
211
   To do so, we start by saving the value of the initial guess x0
212
   (this passage is necessary to avoid overwriting the initial guess during iterations).
213
214
   We then loop over the number of iterations and perform the following steps:
   1) Calculate delta_b from Ax'=b+delta_b where x'=x0+delta_x (the imperfect solution)
216
   2) Perform LU decomposition
   3) Solve Ly = \hat{I}'b for y
218
   4) Solve U Î'x = y for Î'x. This Î'x is used to improve the solution x0: x''=x'-\hat{1}'x (
219
        where x' is the current imperfect solution)
220
221
   # Function for iterative solution using LU decomposition
222
   def iterative_improvement(A, b, x0, iterations):
        Inputs:
        A: matrix
226
        b: vector
227
        x0: initial guess for the solution
228
        iterations: number of iterations to perform
229
        Returns:
230
        x: improved solution
231
```

```
, , ,
232
233
        x = np.copy(x0)
234
        for _ in range(iterations):
             # Difference between actual b and b calculated using current solution x
236
             delta_b = np.dot(A, x) - b
237
238
             # Perform LU decomposition
             L, U = lu_decomposition(A)
239
             \# Solve Ly = \hat{I}'b for y
240
             y = forward_substitution(L, delta_b)
241
             # Solve U \hat{I}'x = y for \hat{I}'x
242
             delta_x = backward_substitution(U, y)
243
244
             # Improve the solution
             x = delta_x
245
        return x
246
247
248
   # Perform 1 LU iteration
   c1 = iterative\_improvement(V, y, c, 1)
   yya1 = np.array([np.sum(c1 * xi ** np.arange(len(c1))) for xi in xx])
251
   ya1 = np.array([np.sum(c1 * xi ** np.arange(len(c1))) for xi in x])
252
   # Perform 10 LU iterations
254
255
   c10 = iterative\_improvement(V, y, c, 10)
    yya10 = np.array([np.sum(c10 * xi ** np.arange(len(c10))) for xi in xx]) \\ ya10 = np.array([np.sum(c10 * xi ** np.arange(len(c10))) for xi in x]) 
256
257
258
   # Plot 2c
259
   fig=plt.figure()
260
   gs=fig.add\_gridspec(2,hspace=0,height\_ratios=[2.0,1.0])
261
   axs=gs.subplots(sharex=True,sharey=False)
262
   axs[0].plot(x,y,marker='o',linewidth=0)
263
   plt.xlim(-1,101)
264
   \mathrm{axs}\left[\,0\,\right].\,\mathrm{set}\,\text{\_ylim}\left(\,-400\,,\!400\right)
265
266
   axs[0].set_ylabel('$y$')
   axs[1].set_ylim(1e-16,1e3)
267
   axs[1].set_ylabel('$|y-y_i|$')
axs[1].set_xlabel('$x$')
268
269
   axs[1].set_yscale('log')
270
   line\ ,= axs\ [\,0\,]\,.\ plot\ (\,xx\,,yya1\,,\, lines\,t\,yle=\,'\,dotted\,\,'\,,\,color=\,'\,red\,\,'\,)
271
   line.set_label('LU with 1 iteration')
272
   axs\,[\,1\,].\,plot\,(\,x\,,abs\,(y-ya1\,)\,\,,lines\,t\,yle=\,,\stackrel{'}{d}otted\,\,,\,color=\,,red\,\,,\,)
273
274
   line,=axs[0].plot(xx,yya10,linestyle='dashdot',color='purple')
275
   line.set_label('LU with 10 iterations')
276
   axs[1].plot(x,abs(y-ya10),linestyle='dashdot',color='purple')
278
   axs[0].legend(frameon=False,loc="lower left")
279
   plt.savefig('./plots/my_vandermonde_sol_2c.png',dpi=600)
280
281
282
   Point d): Exection time comparison
283
284
   To measure the execution time of the three methods, we use the timeit module.
285
   We run each method 100 times (specified by 'number') and save the results to a text
286
        file.
287
    - globals() is used to pass the variables defined in the code to the timeit function
288
    - the first instance of timeit is used to measure the execution time
    - the second instance of timeit is used to display the results obtained in previous
290
        step
291
292
293
   # (a) LU decomposition
294
   lu_time = timeit.timeit("lu_decomposition(V)", globals=globals(), number=100)
295
296
297
   # (b) Neville's algorithm
   neville_time = timeit.timeit("np.array([neville(x, y, xi, 20) for xi in xx])", globals=
298
        globals(), number=100)
```

```
# (c) Iterative improvement with LU decomposition (10 iterations)
iterative_time = timeit.timeit("iterative_improvement(V, y, c, 10)", globals=globals(),
number=100)

# Write results to a text file
with open('timing_results.txt', 'w') as file:
file.write(f"LU Decomposition Time: {lu_time:.6f} seconds\n")
file.write(f"Neville's Algorithm Time: {neville_time:.6f} seconds\n")
file.write(f"Iterative Improvement Time: {iterative_time:.6f} seconds\n")
```

vandermonde.py

For question (a), we are performing a polynomial interpolation on a set of data points using a 19th-degree polynomial, evaluated at ≈ 1000 equally-spaced points. The coefficients of the polynomial are determined by solving a system of linear equations using LU decomposition: they are computed in

```
1.745449629462546071e+01
-1.912658070469094298e+02
2.400858901688944798\,\mathrm{e}{+02}
 -1.102904824261681966e+02
2.688278188786916090e+01
 -4.073856572355006200\,\mathrm{e}{+00}
4.175711801925817701e{-01}
-3.045818340884471134e-02
1.634034746488180218e{-03}
-6.591724714078017072e-05
2.028537534785111589\,\mathrm{e}{-06}
 -4.801909239978146986e-08
8.767406678983276242e-10
 -1.230929039236915523e{-11}
1.315942367469549149e{-13}
-1.051190828510772090e-15
6.072075154696671321e{-18}
 -2.394606872271598716e-20
5.766526114820814136e-23
 -6.394781464844003863e - 26
```

coefficients.txt

. We also plot the absolute difference between the given points y_i and our result y(x), i.e. $|y(x) - y_i|$.

From Fig.1, it seems that the polynomial is fitting the data points quite well for most of the range. This is expected as a polynomial of degree (n) can always fit (n+1) data points exactly. However, towards the end (around x=100), the polynomial shoots up drastically. This is a common issue with high-degree polynomial interpolation known as Runge's phenomenon. It is a form of overfitting where the polynomial oscillates significantly at the boundaries of the data set.

The choice of a 19th-degree polynomial for this data might not be the best. While it fits the given data points well, the extreme behavior at the boundaries suggests that it might not generalize well to other data. A lower-degree polynomial or a different type of function might provide a better fit without the extreme oscillations.

The bottom part of the plot shows the absolute difference between the given y_i values and the calculated y(x) values. This represents the error in the polynomial fit at the data points. Since the polynomial fits the data points exactly, these errors are close to zero, going from order 10^{-14} to 10^{-2} . We can note how the error increases at the boundaries, consistently with the issue presented before.

For question (b), we implement Neville's algorithm (see Fig. 2). It works by recursively evaluating a set of polynomials and combining them to form the final interpolating polynomial. Regarding the interpolation, we are getting the same results as LU decomposition, so they have same efficiency in fitting the given data points. The errors are drastically small for Neville's (reaching a maximum of 10^{-14}) while LU decomposition have errors slightly higher. We could explain this addressing to the nature of the algorithm, which tends to be stable and accurate, especially when the data points are equally spaced. The accuracy of LU decomposition for interpolation can be influenced by the condition of the

matrix involved. Ill-conditioned matrices may lead to numerical instability and higher errors. In our case, the condition number of the Vandermonde matrix depends on the arrangement and spacing of the data points. If the data points are well-spaced and not too close to each other, the condition number may be reasonable. However, as the data points become closely spaced or nearly collinear, the problem of ill-conditioning can arise.

For question (c), we improve iteratively the solution found by LU decomposition. We do such procedure for 1 and 10 iterations (Fig. 3). The interpolation is the same for both cases, since we have overlapping fits. Regarding the trend for the error, the two procedures get similar results, with the LU with 10 iterations having slightly lower errors in some regions of the data.

For question (d), we obtain the execution times for LU decomposition, Neville's algorithm, and LU with iterative improvement: the results are given by

```
LU Decomposition Time: 0.007013 seconds
Neville's Algorithm Time: 9.799903 seconds
Iterative Improvement Time: 0.104062 seconds
```

timing_results.txt

.

- LU Decomposition Time: The algorithm is known for its efficiency in solving linear systems of equations, and it typically has a low time complexity. This outcome indicates that LU decomposition is well-suited for our data set.
- Neville's Algorithm Time: The algorithm involves several polynomial evaluations and recursive calculations: the complexity of these calculations can lead to longer execution times, especially as the number of data points increases.
- Iterative Improvement Time: It is refining a solution through iterative steps, in a relatively fast process. The efficiency of this process can depend on the convergence behavior of the iterative method and the initial guess.

From these results, the LU decomposition method appears to be the most efficient in terms of computational time. If we take the speed as a critical factor and the accuracy achieved by LU decomposition is sufficient for your application, then LU decomposition may be the preferred choice. In our case, the accuracy achieved by LU decomposition is sufficient for our application. However, we have to keep in mind that LU decomposition, while efficient, might be sensitive to ill-conditioned matrices, potentially affecting accuracy in certain cases. In general, accuracy depends on the interpolation requirements and the nature of your data. However, Neville's algorithm is generally considered accurate, especially for well-behaved, evenly spaced data. In our implementation, it is behaving considerably slower than the other, so it might not be the suitable option for this particular setting, even though it is reaching the best performance in terms of lowest errors in the interpolation.

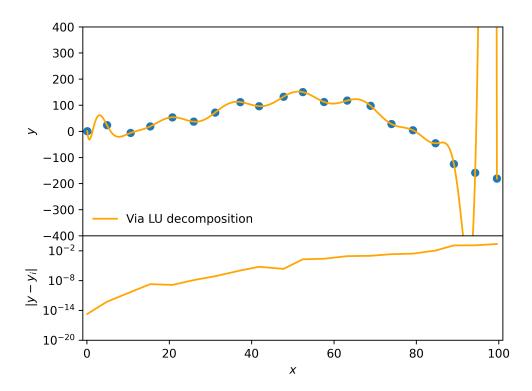


Figure 1: Upper panel: Interpolation on a set of given data points via LU decomposition. The fit is going through all the data points exactly. Nevertheless, it has an oscillating behaviour in correspondence of the last two points. Bottom panel: Absolute difference between the given points y_i and our result y(x), i.e. $|y(x) - y_i|$. The error holds at values close to zero, with a small increase towards the boundaries.

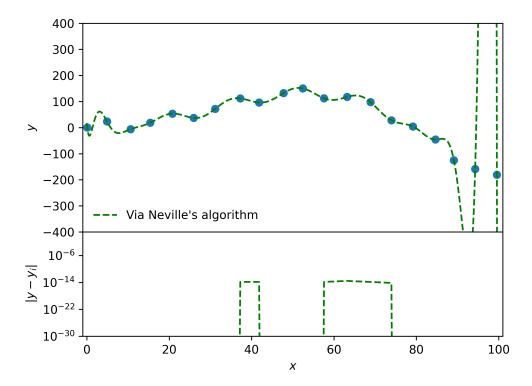


Figure 2: Upper panel: Interpolation on a set of given data points via Neville's algorithm. The polynomial obtained fits the data points well in the range considered, in the same way the interpolation of LU decomposition does. Bottom panel: Absolute difference between the given points y_i and our result y(x), i.e. $|y(x) - y_i|$. The errors of Neville's algorithm reach very small values with respect to LU decomposition's ones.

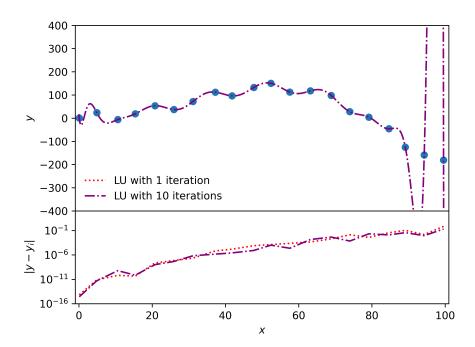


Figure 3: Upper panel: Interpolation on a set of given data points via LU decomposition, with iterative improvement of the found solution (done with 1 and 10 iterations). Both the fits are going through all the data points exactly, hence the overlap. Bottom panel: Absolute difference between the given points y_i and our result y(x), i.e. $|y(x) - y_i|$. The trend is similar for the two implementations, with a small decrease in values reached by LU with 10 iterations.